

Aparapi Quick Reference Guide

Create A Kernel by extending com.amd.aparapi.Kernel

Kernel overrides `run()` method.

Kernel `run()` and `run-reachable` methods can read/write elements of single dimension array primitive final and non-final fields.

Kernel can read primitive final and non-final fields

Create an anonymous inner class extending com.amd.aparapi.Kernel

Kernel must override `run()` method.

Kernel `run()` and `run-reachable` methods can read and write elements of captured single dimension arrays of primitive from the call-site

Kernel can read captured final primitive values from the call-site.

This code sets each element of `data[]` to its index.

Executing your kernel over a given range (0..<range>)

Use `Kernel.execute(int <range>)` to execute over the range `0..<range>`

Each Kernel execution will receive a unique value (`0..<range>`) via `kernel.getGlobalId()`.

This code each element of `data[]` to its index.

Compiling your application

Add `aparapi.jar` to your classpath

Turn on debugging information for compiled classes (`javac -g` option)

Linux® users use `:` as a path separator instead of `;`

Running your application

Add `aparapi.jar` to your classpath

Add `aparapi` shared library to `java.library.path`

Modes of execution

Mode	Method Of Execution	OpenCL™
JTP	Java Thread Pool. One Thread per available core	
SEQ	Single sequential Java loop.	
GPU	Convert bytecode to OpenCL execute on first available GPU device	✓
CPU	Convert bytecode to OpenCL execute on available CPU cores.	✓

Default to GPU if OpenCL is available

Falls back to JTP mode if Kernel cannot be converted or OpenCL unavailable.

Setting default execution mode

Developer request mode prior to first call to `kernel.execute(<range>)` using `kernel.setExecutionMode()`.

Alternatively set the requested mode (for all Kernels) via `com.amd.aparapi.ExecutionMode` property

Determining the execution mode

Determine the actual execution mode after `Kernel.execute(<range>)` returns via `Kernel.getExecutionMode()`.


Alternatively set `com.amd.aparapi.enableReportKernelExecutionMode` property to true when application is launched.

Aparapi properties to control runtime behavior

Properties	Options (default)
<code>com.amd.aparapi.executionMode</code>	SEQ JTP CPU GPU
<code>com.amd.aparapi.enableExecutionModeReporting</code>	true false
<code>com.amd.aparapi.enableShowGeneratedOpenCL</code>	true false
<code>com.amd.aparapi.logLevel</code>	FINEST FINE INFO WARNING SEVERE
<code>com.amd.aparapi.enableProfiling</code>	true false
<code>com.amd.aparapi.enableVerboseJNI</code>	true false
<code>com.amd.aparapi.instructionListenerClass</code>	Name of a class implements <code>Config.InstructionListener</code>

For example:

What happens when we call Kernel.execute(<range>)?



Explicit buffer management

For some applications where Kernel.execute calls are executed multiple times (loops) and where the Java code between Kernel.execute calls does not access buffers (primitive arrays) Aparapi can cause unnecessary buffer copies.

For performance reasons we allow the developer to take explicit control of transfers.

Kernel.execute(range), Kernel.put() and Kernel.get() support the fluent style of API (return Kernel instance) so that calls can be chained.

Kernel methods to determine execution identity


Kernel.execute(range, passCount) will execute the Kernel.run() method 0..<range> for each iteration through an outer the loop 1..<passCount>.

Note Kernel.execute(<range>) is equivalent to Kernel.execute(<range>, 1)

The 0..<range> execution is divided into one or more equal sized groups.

For example; Kernel.execute(16) may result in 1 group of 16, 2 groups of 8, 4 groups of 4 or 16 groups of 1

For example if Kernel.execute(16) resulted in 2 groups of 8 the following diagram shows typical values for each corresponding global id value we might expected for each of the kernel identity methods.



Choose <range> carefully (powers of 2 ideal) to maximize group size.

Note if <range> is a prime number then group size will be 1 and inefficient

Method	Returns
getGlobalId()	The global id that this Kernel execution represents (0..<range>)
getGroupSize()	The size of each group. The same value returned for all Kernel executions resulting from one invocation of Kernel.execute(<range>).
getGroupId()	The groupId of the currently executing Kernel (0..getNumGroups())
getLocalId()	The localId (from 0..<groupSize>) for this Kernel execution.
getGlobalSize()	The value of <range>
getPassId()	Determine the current passId (0..<passCount>).

There is no way to determine group size or count prior to executing Kernel.execute(<range>) and no guarantee that the same values will be used for subsequent invocations of Kernel.execute(<range>).

Extended forms of Kernel.execute()

If Kernel.execute() is sole statement in a loop whose iteration count is known on first iteration we can replace the loop with an extra 'pass count' arg to Kernel.execute(range, passCount).

Can be replaced with

Code is more compact and avoids unnecessary buffer transfers between iterations.

Kernel implementation can extract the value of the current pass using Kernel.getPassId().

Mapping of Aparapi Kernel math methods to Java and OpenCL equivalents

Kernel method	Java mapping	OpenCL mapping
abs(float)	Math.abs(float)	fabs(float)
abs(double)	Math.abs(double)	fabs(double)
abs(int)	Math.abs(int)	abs(int)
abs(long)	Math.abs(long)	abs(long)
acos(float)	Math.acos(float)	acos(float)
acos(double)	Math.acos(double)	acos(double)
asin(float)	Math.asin(float)	asin(float)
asin(double)	Math.asin(double)	asin(double)
atan(float)	Math.atan(float)	atan(float)
atan(double)	Math.atan(double)	atan(double)
atan2(float, float)	Math.atan2(float, float)	atan2(float, float)
atan2(double, double)	Math.atan2(double, double)	atan2(double, double)
ceil(float)	Math.ceil(float)	ceil(float)
ceil(double)	Math.ceil(double)	ceil(double)
cos(float)	Math.cos(float)	cos(float)
cos(double)	Math.cos(double)	cos(double)
exp(float)	Math.exp(float)	exp(float)
exp(double)	Math.exp(double)	exp(double)
floor(float)	Math.floor(float)	floor(float)
floor(double)	Math.floor(double)	floor(double)
max(float, float)	Math.max(float, float)	fmax(float, float)
max(double, double)	Math.max(double, double)	fmax(double, double)
max(int, int)	Math.max(int, int)	max(int, int)
max(long, long)	Math.max(long, long)	max(long, long)
min(float, float)	Math.min(float, float)	fmin(float, float)
min(double, double)	Math.min(double, double)	fmin(double, double)
min(int, int)	Math.min(int, int)	min(int, int)
min(long, long)	Math.min(long, long)	min(long, long)
log(float)	Math.log(float)	log(float)
log(double)	Math.log(double)	log(double)
native_sqrt(float)	See method implementation	native_sqrt(float)
native_sqrt(double)		native_sqrt(double)
pow(float, float)	Math.pow(float, float)	pow(float, float)
pow(double, double)	Math.pow(double, double)	pow(double, double)
IEEERemainder(float, float)	Math.IEEERemainder(float, float)	remainder(float, float)
IEEERemainder(double, double)	Math.IEEERemainder(double, double)	remainder(double, double)
rint(float)	Math.rint(float)	rint(float)
rint(double)	Math.rint(double)	rint(double)
round(float)	Math.round(float)	round(float)
round(double)	Math.round(double)	round(double)
rsqrt(float)	1f/Math.sqrt(float)	rsqrt(float)
rsqrt(double)	1.0/Math.sqrt(double)	rsqrt(double)
sin(float)	Math.sin(float)	sin(float)
sin(double)	Math.sin(double)	sin(double)
sqrt(float)	Math.sqrt(float)	sqrt(float)
sqrt(double)	Math.sqrt(double)	sqrt(double)
tan(float)	Math.tan(float)	tan(float)
tan(double)	Math.tan(double)	tan(double)
toRadians(float)	Math.toRadians(float)	radians(float)
toRadians(double)	Math.toRadians(double)	radians(double)
toDegrees(float)	Math.toDegrees(float)	degrees(float)
toDegrees(double)	Math.toDegrees(double)	degrees(double)

Note the differences in precision between Java and OpenCL™ implementation of arithmetic functions to determine whether the difference in precision is acceptable.

OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.